

Pintos Task 3: Virtual Memory

Eva Kalyvianaki
Mark Rutland

Imperial College London

24 February 2011

Introduction

Goal

Remove the space limitations of main memory

Tasks

- i Lazy-loading of executables: load pages on demand
- ii Stack growth
- iii Memory-mapped files
- iv Swapping

Getting Started

Does Task 2 need to be complete?

- Completion of Task 2 is essential for Task 3
- Virtual memory is an extension to user program handling
- A good understanding of threads and processes in pintos is required

Testing

- tests for Task 3
 - Including Task 2 tests
- Tests take a long time to run
 - Running individual tests will allow you to work faster

Important Files

Files to modify:

- i threads/thread.h
 - Add supplemental page table to struct thread
- ii userprog/process.c
 - Handles loading, execution and exit of processes
 - Handles stack setup
- iii userprog/syscall.c
 - Handles all syscalls
 - Checks validity of process-supplied data
- iv userprog/exception.c
 - Handles segmentation faults
 - Some lazy-loading code goes here

Warning

Some features you are to add should not be placed in existing files, as they will require a large amount of additional code.

Files to create:

- i Frame table
 - `vm/frame.c` ?
- ii Supplemental page table
 - `vm/page.c` ?
- iii Swap table
 - `vm/swap.c` ?
- iv Memory map table
 - `vm/mmap.c` ?

Important Files

Files with useful functions

- i `<bitmap.h>`
 - In-kernel bitmap implementation
- ii `<hash.h>`
 - Generic hash table implementation
- iii `"devices/block.c"`
 - Provides sector-based access to block devices
- iv `"threads/vaddr.h"`
 - `is_user_vaddr()` used to validate user provided pointers
 - `pg_ofs()` used to find the offset of a pointer within a page
 - `pg_round_down()` rounds a pointer down to the nearest page boundary
- v `"userprog/pagedir.c"`
 - Functions for querying and altering the page table

Virtual Memory in Pintos

Kernel addresses

- Due to the way x86 virtual memory is implemented, it is not possible to address memory by physical address whilst in protected mode
- You will not need to use physical addresses directly, kernel functions take **kernel addresses**
- Pintos maps every physical address to a **kernel address** using a fixed offset (`PHYS_BASE`)

Page Faults in Pintos

- Handled in `page_fault` in `userprog/exception.c`
- Currently not very useful:
 - Kills user processes on any page fault

Project Requirements

- Frame allocator
- Lazy-loading of executables
- Dynamic stack allocation
- Memory mapped files
- Frame/swap reclamation on exit
- Swapping
- Eviction policy

Frame Allocator + Eviction Policy

What?

- After a number of allocations, physical frames will be exhausted
- Need a way of tracking allocations to enable swapping

How?

- Build a frame table
- Contains pointers to pages and other data of your choice
- Record which thread owns a frame
- Most important operation is to obtain an used frame
- Modify calls to `palloc_get_page()` and `palloc_free_page()`

Warning

Many processes may need to allocate or free pages simultaneously. Carefully consider how you will handle concurrent access.

Segment setup

Current process setup

- Executable is loaded into process memory by `load_segment()`
- Stack segment is ignored
- Stack is set up at a fixed address

So...?

- The entire executable must be loaded before it is run
- Data which may never be used is read into memory

Solution?

- Use the supplemental page table
- Load in segments on demand, using the page-fault handler

Stack growth

The problem

- Allocating a small stack (as Pintos currently does) prohibits use of large stack objects
- Allocating a large stack wastes space for processes which do not use it

The solution

- Grow the stack dynamically, as page faults occur
- Identify stack accesses by looking at the `esp` of the struct `intr_frame`

Stack accesses

- As with any other page, the stack can be swapped out.
- Set an absolute limit on stack size, 8 MB.

Memory-mapped files

Idea

- Map files into process address space
- On page fault, load file data into page
- When swapping, write file pages back to disk

Mapping files

- Need to connect mapids to actual files (a table?)
- Multiple processes can map a single file (share the page?)

Warning(s)

- Maps should remain if a file is closed or removed (use `file_reopen`)
- All mapped files should be flushed when a process exits